

Analyzing Fashion Trends and Customer Preferences Using Big Data Technologies

Group 6

Avirit Singh

Jay Joshi

Tanu Datt

Vaibhavi Rao

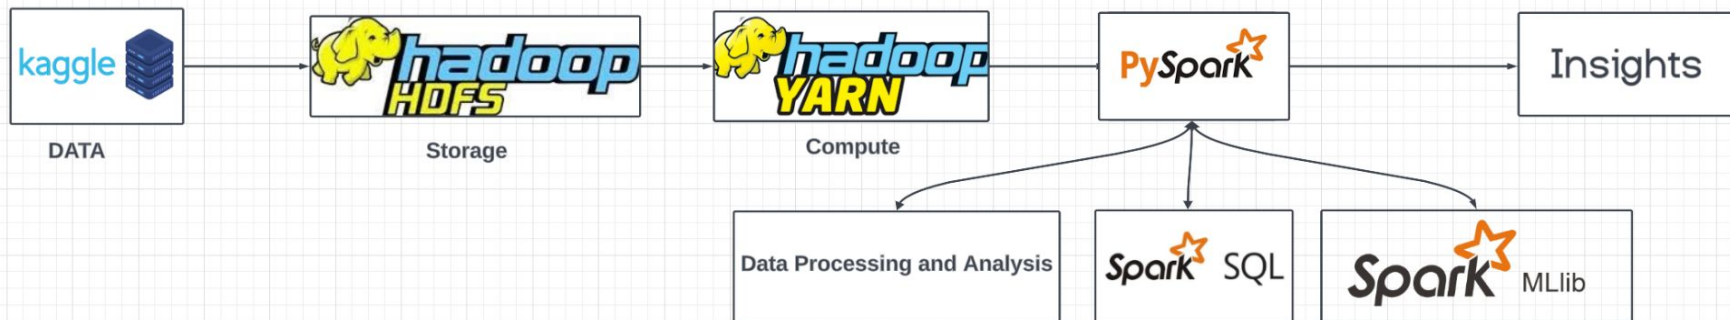
Varun Patil

Pragya Priyadarshini

Problem Statement

The most important challenge for fashion retailers is having to process efficiently and analyze large, varied data in order to arrive at actionable insights for strategic decisions. This project addresses the requirement for advanced analytic models capable of processing such huge quantities of fashion data to predict future trends and consumer preferences with accuracy. Big data technology coupled with machine learning is applied herein to enhance inventory optimization and marketing effectiveness with a view toward improvement in profitability and customer satisfaction for apparel retailers.

Proposed Method



Dataset Overview

Structure

- **Directories**
 - **images/:** Contains 44,400 high-resolution product images (JPG format), each mapped to a unique product ID.
 - **styles/:** Includes metadata files (JSON format) with detailed product descriptions and attributes.
- **Files**
 - **styles.csv:** Maps product IDs to 12 key attributes (e.g., categories, display names, labels).
 - **images.csv:** Provides supplementary information about the image data.

Content

- Each product is linked through a unique product ID, connecting images, metadata, and CSV records.
- **styles.csv** serves as the primary mapping file, detailing product attributes.
- JSON metadata complements visual data with rich product context.

Summary

- **Total Files:** ~88,900
- **Directories:** 2 (images, styles)
- **File Formats:** JPG, JSON, CSV
- **Dataset Size:** ~23.1 GB

Dataset Pre-Processing

Efficient handling of JSON

Defined a schema to process nested structures, ensuring coherence in data types.

Flattened complex fields like meta and data to query them more easily- for example, meta.code, data.price.

Exploded the arrays in individual rows to further develop the relational analysis that could be done with data_crossLinks.

Data Cleaning and Validation

Addressed null/NaN values in critical columns, such as label, via checks, filters, and replacements.

Dropped irrelevant or redundant columns: notification, data_styleImages etc.

Ensured completeness and consistency in the final data for analysis.

Feature Selection and Transformations

Focused on main attributes such as data_id, data_price, data_brandName, and meta_code. Extracted new features, for example discount calculation: data.price - data.discountedPrice. Prepared data for EDA, feature engineering, and machine learning workflows.

Data Analysis[Insights]

A./ Decoding Sales Performance: A Multifaceted Data Analysis Approach

Top 10 Brands by Total Revenue

```
from pyspark.sql import functions as F

# Calculate total revenue (price * quantity) for each brand
top_brands_revenue = final_df.groupBy("data_brandName") \
    .agg(F.sum(F.col("data_discountedPrice")).alias("total_revenue")) \
    .orderBy(F.desc("total_revenue"))

top_brands_revenue.show(10, truncate=False)
```

data_brandName	total_revenue
Nike	5918015
Puma	4418246
ADIDAS	3955147
United Colors of Benetton	2678491
Fossil	1698697
CASIO	1355145
FNF	1162341
French Connection	1142081
Catwalk	1090089
Timberland	1017409

only showing top 10 rows

Data Analysis[Insights]

Finding the Highest Priced Nike (any brand)Product

```
from pyspark.sql import functions as F
import matplotlib.pyplot as plt
from PIL import Image
import os

brand = input(str("Enter your Brand: "))
# Step 1: Filter rows where brand is "Nike"
nike_data = cleaned_df.filter(cleaned_df['data_brandName'] == brand)

# Step 2: Find the row with the highest data_price
# Use PySpark to find the row with the highest price
max_price_row = nike_data.orderBy(F.desc("data_price")).limit(1).collect()[0]

# Extract values from the row
highest_price_id = max_price_row['data_id']
data_price = max_price_row['data_price']
data_productDisplayName = max_price_row['data_productDisplayName']
data_brandName = max_price_row['data_brandName']
data_ageGroup = max_price_row['data_ageGroup']
data_gender = max_price_row['data_gender']
data_displayCategories = max_price_row['data_displayCategories']

# Print the respective values
print(f"The data_id with the highest data_price for {brand} is: INR {highest_price_id}")
print(f"Details of the product with the highest price:")
print(f"Price: INR {data_price}")
print(f"Product Display Name: {data_productDisplayName}")
print(f"Brand Name: {data_brandName}")
print(f"Age Group: {data_ageGroup}")
print(f"Gender: {data_gender}")
print(f"Display Categories: {data_displayCategories}")
```

Enter your Brand: Nike

[Stage 188:=====> (1350 + 8) / 1389]

The data_id with the highest data_price for Nike is: INR 44235

Details of the product with the highest price:

Price: INR 12995

Product Display Name: Nike Men Air Max+ 2012 Blue Sports Shoes

Brand Name: Nike

Age Group: Adults-Men

Gender: Men

Display Categories: Footwear

Data Analysis[Insights]

Profitability Metrics by Category

```
# Calculate profitability by category
category_profitability = final_df.groupby("data_displayCategories").agg(
    avg("data_vat").alias("avg_vat"),
    count("*").alias("sales_volume"),
    sum("data_discountedPrice").alias("total_revenue")
).orderBy(col("total_revenue").desc())

# Show results
category_profitability.show(truncate=False)
```

data_displayCategories	avg_vat	sales_volume	total_revenue
Accessories	13.592643194955334	9515	21463790
Footwear	14.498714836498644	7003	16953850
Casual Wear	5.501028101439342	8754	9572107
NULL	9.952228749136143	5788	7395242
Footwear,Sale	14.389261744966444	894	2430895
Ethnic Wear	5.5	2082	1886113
Casual Wear,Sale	5.5	1489	1518705
Casual Wear,Winterwear	5.5	649	1281982
Sports Wear	5.540723981900452	884	1214576
Formal Wear	5.5	1012	1191554
Accessories,Sale	12.941558441558442	462	589937
Innerwear	8.025538461538462	1625	555310
Sports Shoes,Footwear and Clearance,Sale and Clearance,Footwear,Sale	14.5	143	509620
Sports Shoes,Footwear	14.5	99	374836
Sale and Clearance,Footwear,Sale	14.409090909090908	99	346462
Sports Wear,Winterwear	5.589108910891089	101	324167
Sports Wear,Sale	5.544334975369458	203	290207
Tshirts,Casual Wear and Clearance,Sale and Clearance,Casual Wear,Sale	5.5	384	259297
Footwear and Clearance,Sale and Clearance,Footwear,Sale	14.5	95	221314
Shirts,Casual Wear	5.5	134	206887

only showing top 20 rows

Data Analysis[Insights]

Top Brands by Usage Segment and Gender

```
top_brands_segment = final_df.groupBy("data_brandName", "data_usage", "data_gender").agg(  
    count("*").alias("usage_count")  
).orderBy("usage_count", ascending=False)  
  
# Show results  
top_brands_segment.show(truncate=False)
```

data_brandName	data_usage	data_gender	usage_count
Nike	Sports	Men	1030
Puma	Casual	Men	1008
United Colors of Benetton	Casual	Men	792
Catwalk	Casual	Women	732
ADIDAS	Casual	Men	688
United Colors of Benetton	Casual	Women	679
ADIDAS	Sports	Men	666
Baggit	Casual	Women	625
Fabindia	Ethnic	Women	550
Lino Perros	Casual	Women	497
Nike	Casual	Men	455
Wrangler	Casual	Men	442
Puma	Sports	Men	412
Jealous 21	Casual	Women	402
Murcia	Casual	Women	370
Colorbar	Casual	Women	357
Myntra	Casual	Men	352
Nike	Sports	Women	338
W	Ethnic	Women	337
Femella	Casual	Women	334

only showing top 20 rows

Data Analysis[Insights]

Best-Selling Categories by Gender

```
# Best-selling categories by gender
final_df.groupBy("data_gender", "data_displayCategories") \
    .count() \
    .orderBy("count", ascending=False) \
    .show(20)
```

data_gender	data_displayCategories	count
Men	Casual Wear	4859
Men	Footwear	4559
Women	NULL	4457
Men	Accessories	4215
Women	Accessories	4175
Women	Casual Wear	2954
Women	Ethnic Wear	1986
Women	Footwear	1951
Unisex	Accessories	1087
Men	NULL	1011
Men	Formal Wear	953
Men	Casual Wear,Sale	951
Men	Innerwear	833
Women	Innerwear	783
Men	Sports Wear	708
Men	Footwear,Sale	591
Boys	Casual Wear	548
Women	Casual Wear,Sale	479
Men	Casual Wear,Winte...	475
Girls	Casual Wear	384

only showing top 20 rows

Data Analysis[Insights]

Total Revenue by Base Color [popularity]

```
# Group by base color and calculate total revenue
color_revenue = final_df.groupby("data_baseColour").agg(
    sum("data_discountedPrice").alias("total_revenue")
).orderBy("total_revenue", ascending=False)

# Show the results
color_revenue.show(truncate=False)
```

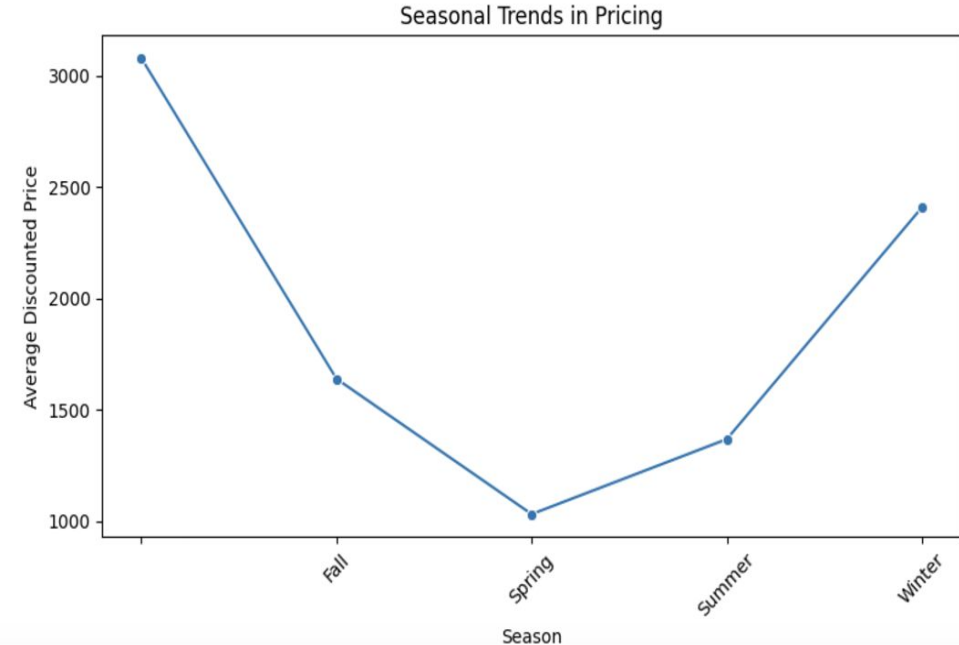
data_baseColour	total_revenue
Black	19197911
White	10425403
Blue	6585835
Brown	6262574
Grey	4572663
Silver	3215362
Red	2998091
Green	2422765
Navy Blue	2233161
Purple	1988349
Pink	1861297
Gold	1200152
Beige	1016312
Steel	987989
Yellow	804199
Maroon	672562
Orange	656709
Olive	650430
Cream	577971
Multi	563600

only showing top 20 rows

Data Analysis[Insights]

B.] Seasonal Trends in Discounted Pricing

Seasonal Trends in Average Discounted Pricing



```
from pyspark.sql import functions as F
import matplotlib.pyplot as plt
import seaborn as sns

# Calculate the average discounted price by season
seasonal_price = final_df.groupBy("data_season") \
    .agg(F.avg("data_discountedPrice").alias("avg_discounted_price")) \
    .orderBy("data_season")

# Convert to Pandas for visualization
seasonal_price_pd = seasonal_price.toPandas()

# Visualization: Line graph
plt.figure(figsize=(8, 5))
sns.lineplot(data=seasonal_price_pd, x="data_season", y="avg_discounted_price", marker="o")
plt.title("Seasonal Trends in Pricing")
plt.ylabel("Average Discounted Price")
plt.xlabel("Season")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

Data Analysis[Insights]

Usage Analysis by Season and Category

```
from pyspark.sql.functions import count

# Group by season and usage
season_usage_analysis = final_df.groupBy("data_season", "data_usage").agg(
    count("*").alias("usage_count")
)

# Show the result
season_usage_analysis.orderBy("data_season", "usage_count", ascending=False).show(truncate=False)
```

data_season	data_usage	usage_count
Winter	Casual	7986
Winter	Formal	247
Winter	Sports	117
Winter	Ethnic	105
Winter	Smart Casual	32
Winter	Travel	12
Winter	NA	5
Winter	Party	5
Winter		1
Summer	Casual	16312
Summer	Sports	2135
Summer	Ethnic	1887
Summer	Formal	1073
Summer	NA	30
Summer	Travel	12
Summer	Smart Casual	10
Summer	Party	10
Spring	Casual	2554
Spring	NA	270
Spring	Sports	110

only showing top 20 rows

Data Analysis[Insights]

Top Products by Revenue and Season

```
from pyspark.sql.functions import col, sum

# Calculate total revenue for each product by season
top_products_season = final_df.groupBy("data_season", "data_productDisplayName").agg(
    sum("data_discountedPrice").alias("total_revenue")
).orderBy("data_season", "total_revenue", ascending=False)

# Show the top products by season
top_products_season.show(truncate=False)
```

data_season	data_productDisplayName	total_revenue
Winter	Nautica Men Black Dial Chronograph Watch	156435
Winter	Timex Men Black Dial Watch	120140
Winter	Titan Men White Dial Watch	105200
Winter	Giordano Men Black Dial Watch	102528
Winter	Morellato Men Silver Dial Watch	83250
Winter	Fastrack Men Black Dial Watch	78745
Winter	Ed Hardy Men Black Dial Watch	74980
Winter	Citizen Men Black Dial Chronograph Watch	73200
Winter	Titan Men Black Dial Watch	72930
Winter	Catwalk Women Black Heels	70485
Winter	Miss Sixty Silver Dial Watch	68660
Winter	Ray-Ban Men Aviator Sunglasses	68616
Winter	Titan Men Black Watch	68110
Winter	Red Tape Men Brown Shoes	64370
Winter	Citizen Men Black Dial Eco-Drive Watch	63600
Winter	Maxima Men White Dial Watch	63574
Winter	Giordano Men White Dial Watch	62396
Winter	Ray-Ban Men Aviator Gold Sunglasses	60723
Winter	Citizen Women White Dial Watch	59700
Winter	United Colors of Benetton Men Black Sunglasses	58005

only showing top 20 rows

Data Analysis[Insights]

Seasonal and Yearly Distribution of Base Colors in Products

data_season	data_year	data_baseColour	count
Fall	2012	Navy Blue	64
Winter	2017	Black	3
Summer	2012	Navy Blue	586
Summer	2012	Purple	474
Fall	2018	White	3
Fall	2011	Blue	1338
Fall	2011	White	1148
Fall	2011	Grey	760
Fall	2012	Charcoal	3
Summer	2017	Navy Blue	9
Summer	2013	Black	163
Summer	2015	Red	21
Winter	2012	Blue	73
Winter	2015	Black	553
Winter	2012	Coffee Brown	3
Winter	2011	Yellow	3
Spring	2013	Red	14
Winter	2012	Brown	168
Spring	2013	Black	92
Summer	2014	Black	34

only showing top 20 rows

Data Analysis[Insights]

Multi-Level Revenue and Sales Count by Brand, Style, and Season

```
from pyspark.sql.functions import col, sum, count

# Calculate revenue at multiple levels
multi_level_revenue = final_df.groupBy("data_brandName", "data_styleType", "data_season").agg(
    sum("data_discountedPrice").alias("total_revenue"),
    count("*").alias("sales_count")
).orderBy("data_brandName", "total_revenue", ascending=False)

# Show the results
multi_level_revenue.show(truncate=False)
```

data_brandName	data_styleType	data_season	total_revenue	sales_count
yellowe	P	Summer	51480	32
yellowe	P	Fall	790	1
vogue	P	Winter	76743	16
test	D	Spring	500	1
s.Oliver	P	Fall	136233	67
s.Oliver	P	Summer	93139	61
roxy	P	Winter	109488	59
roxy	P	Summer	2390	2
pierre cardin	P	Spring	3050	2
maxima	P	Winter	414041	252
maxima	RTV	Winter	11115	4
maxima	P	Summer	4415	2
maxima	DEL	Winter	1174	2
ice watch	P	Winter	102915	17
iPanema	P	Winter	75737	93
iPanema	P	Summer	999	1
iPanema	CDL	Winter	699	1
dunhill	P	Spring	78450	24
aramis	P	Spring	8265	3
Yves Saint Laurent	P	Spring	8800	2

only showing top 20 rows

Data Analysis[Insights]

C./ Retail Insights: Discount, Customer, and Pricing Analysis

Customer Segmentation by Spending Levels

```
# Group customers by total spending
customer_segments = final_df.groupBy("data_id").agg(
    sum("data_discountedPrice").alias("total_spent")
).withColumn(
    "spending_segment",
    when(col("total_spent") > 10000, "High Spender")
    .when(col("total_spent") > 5000, "Medium Spender")
    .otherwise("Low Spender")
).groupBy("spending_segment").agg(
    count("*").alias("customer_count"),
    avg("total_spent").alias("avg_spent")
)

# Show the results
customer_segments.show(truncate=False)
```

spending_segment	customer_count	avg_spent
Medium Spender	1764	7012.326530612245
High Spender	203	12590.536945812808
Low Spender	42479	1338.4696074044418

Data Analysis[Insights]

Potential Revenue Loss

```
from pyspark.sql.functions import col, count, sum

# Filter for unsold winter products from previous years
unsold_products = final_df.filter(
    (col("data_year") < 2023) & (col("data_season") == "Winter")
).groupBy("data_productDisplayName", "data_brandName").agg(
    count("*").alias("unsold_count"),
    sum("label").alias("potential_revenue_loss") # Assuming 'label' is the original price
).orderBy(col("potential_revenue_loss").desc())

# Show results
unsold_products.show(truncate=False)
```

data_productDisplayName	data_brandName	unsold_count	potential_revenue_loss
Nautica Men Black Dial Chronograph Watch	Nautica	13	156435
Timex Men Black Dial Watch	Timex	22	120140
Giordano Men Black Dial Watch	GIORDANO	19	110450
Titan Men White Dial Watch	Titan	26	105200
Morellato Men Silver Dial Watch	Morellato	9	83250
Fastrack Men Black Dial Watch	Fastrack	41	78745
Giordano Men White Dial Watch	GIORDANO	14	76400
Ray-Ban Men Aviator Sunglasses	Ray-Ban	16	76240
Ed Hardy Men Black Dial Watch	Ed Hardy	8	74980
Citizen Men Black Dial Chronograph Watch	Citizen	7	73200
Titan Men Black Dial Watch	Titan	15	72930
Catwalk Women Black Heels	Catwalk	43	70485
Miss Sixty Silver Dial Watch	MISS SIXTY	8	68660
Titan Men Black Watch	Titan	16	68110
Ray-Ban Men Aviator Gold Sunglasses	Ray-Ban	13	67470
United Colors of Benetton Men Black Sunglasses	United Colors of Benetton	18	64450
Red Tape Men Brown Shoes	Red Tape	26	64370
Polaroid Men Sunglasses	Polaroid	18	64291
Citizen Men Black Dial Eco-Drive Watch	Citizen	6	63600
Maxima Men White Dial Watch	maxima	37	63574

only showing top 20 rows

Data Analysis[Insights]

Distinct Product purchased by each customer

```
from pyspark.sql.functions import countDistinct

# Count distinct products purchased by each customer
customer_loyalty = final_df.groupBy("data_id", "data_gender", "data_ageGroup").agg(
    countDistinct("data_productDisplayName").alias("unique_purchases"),
    sum("data_discountedPrice").alias("total_spent")
).orderBy("total_spent", ascending=False)

# Show the results
customer_loyalty.show(truncate=False)
```

data_id	data_gender	data_ageGroup	unique_purchases	total_spent
35288	Unisex	Adults-Unisex	1	28950
35282	Unisex	Adults-Unisex	1	21950
52686	Men	Adults-Men	1	21220
59253	Men	Adults-Men	1	18995
5062	Men	Adults-Men	1	18900
28438	Women	Adults-Women	1	17995
53014	Women	Adults-Women	1	17500
52691	Women	Adults-Women	1	17150
52690	Women	Adults-Women	1	17000
29945	Men	Adults-Men	1	16450
29923	Men	Adults-Men	1	15550
51639	Men	Adults-Men	1	15495
53020	Women	Adults-Women	1	15426
29950	Men	Adults-Men	1	15350
29951	Men	Adults-Men	1	15350
29937	Men	Adults-Men	1	15350
29934	Women	Adults-Women	1	15350
29944	Men	Adults-Men	1	15350
53016	Women	Adults-Women	1	15150
52688	Women	Adults-Women	1	15050

only showing top 20 rows

Data Analysis[Insights]

Pricing Analysis by Gender, Season, and Price Range

```
from pyspark.sql.functions import when

# Group products into price ranges
pricing_analysis = final_df.withColumn(
    "price_range",
    when(col("data_discountedPrice") < 500, "<500")
    .when((col("data_discountedPrice") >= 500) & (col("data_discountedPrice") < 1500), "500-1500")
    .when((col("data_discountedPrice") >= 1500) & (col("data_discountedPrice") < 3000), "1500-3000")
    .otherwise(">3000")
).groupBy("data_gender", "data_season", "price_range").agg(
    count("*").alias("sales_count"),
    sum("data_discountedPrice").alias("total_revenue")
).orderBy("data_gender", "data_season", col("sales_count").desc())

# Show results
pricing_analysis.show(truncate=False)
```

data_gender	data_season	price_range	sales_count	total_revenue
Boys	Fall	<500	78	28957
Boys	Fall	500-1500	27	23091
Boys	Fall	>3000	5	20850
Boys	Fall	1500-3000	3	6497
Boys	Spring	<500	5	1618
Boys	Spring	500-1500	1	999
Boys	Summer	<500	551	160526
Boys	Summer	500-1500	137	111688
Boys	Summer	1500-3000	2	3398
Boys	Summer	>3000	1	3295
Boys	Winter	<500	18	5548
Boys	Winter	500-1500	2	1898
Girls	Fall	<500	60	23099
Girls	Fall	500-1500	21	17863
Girls	Spring	<500	3	1357
Girls	Summer	<500	419	124950
Girls	Summer	500-1500	114	92938
Girls	Summer	1500-3000	3	5097
Girls	Winter	<500	22	6894
Girls	Winter	500-1500	13	9372

only showing top 20 rows

Data Analysis[Insights]

Discount Analysis by Range, Sales Count, and Total Revenue

```
from pyspark.sql.functions import when, col

# Calculate optimal discount range
discount_analysis = final_df.withColumn(
    "discount_percentage",
    ((col("label") - col("data_discountedPrice")) / col("label")) * 100
).withColumn(
    "discount_range",
    when(col("discount_percentage") < 10, "<10%")
    .when((col("discount_percentage") >= 10) & (col("discount_percentage") < 30), "10-30%")
    .when((col("discount_percentage") >= 30) & (col("discount_percentage") < 50), "30-50%")
    .otherwise(">50%")
).groupBy("discount_range").agg(
    count("*").alias("sales_count"),
    sum("data_discountedPrice").alias("total_revenue")
).orderBy(col("sales_count").desc())

# Show results
discount_analysis.show(truncate=False)
```

discount_range	sales_count	total_revenue
<10%	37300	65053034
10-30%	3056	3835260
>50%	2916	1544326
30-50%	1174	1325761

ML Model

```
from pyspark.ml.evaluation import RegressionEvaluator

# Evaluate RMSE
evaluator = RegressionEvaluator(labelCol="label", predictionCol="prediction", metricName="rmse")
rmse = evaluator.evaluate(lr_predictions)
print(f"Root Mean Square Error (RMSE) Linear Regression: {rmse}")

# Evaluate R2
r2_evaluator = RegressionEvaluator(labelCol="label", predictionCol="prediction", metricName="r2")
r2 = r2_evaluator.evaluate(lr_predictions)
print(f"R2 Linear Regression: {r2}")
```

Root Mean Square Error (RMSE) Linear Regression: 0.0999982838823679

[Stage 44:=====> (24 + 8) / 36]

R² Linear Regression: 0.999999996604294

```
# Print coefficients and intercept
print(f"Coefficients LR: {lr_model.coefficients}")
print(f"Intercept LR: {lr_model.intercept}")
```

Coefficients LR: [0.0,0.0,0.0,0.9999417273126453,0.0]

Intercept LR: 0.09854435853073856

ML Model

Descriptive Features:

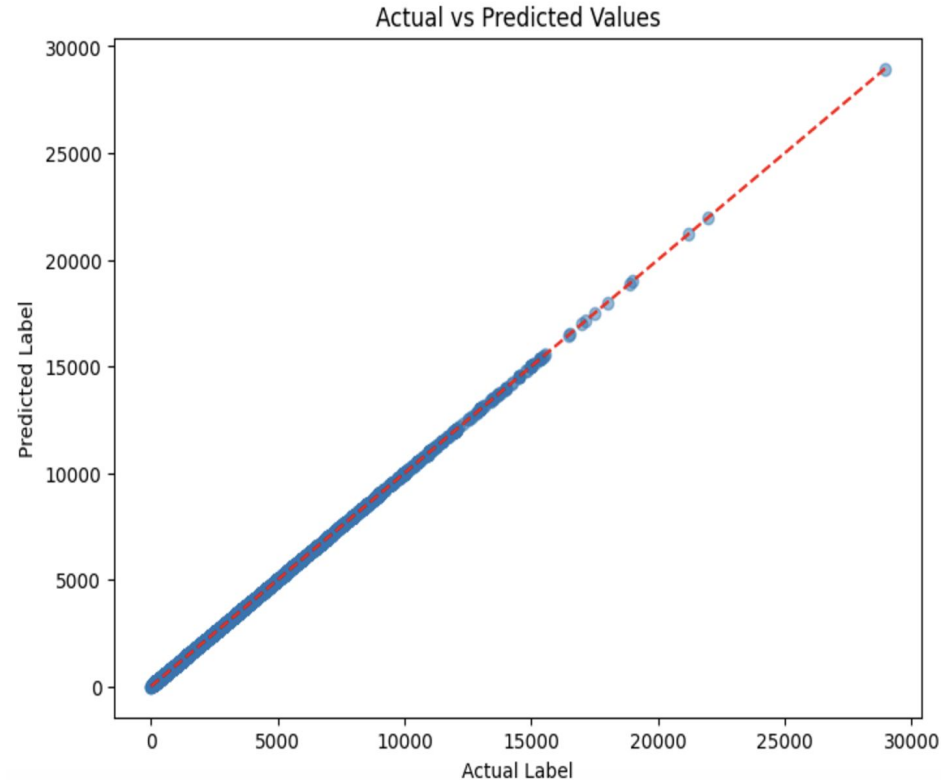
- Categorical Features: Indexed (data_fashionType, data_brandName, data_baseColour)
- Numerical Features: Used directly (data_price, data_year)

Target Feature: data_price

Model Implemented: linear regression

Model Performance:

- **R² Value:** 0.09 - Indicates perfect prediction accuracy.
- **RMSE:** Remarkably low - Signifies the model's predictions precisely match the actual values.
- **Visual:** Small plot showing Actual vs. Predicted Values illustrating the direct correspondence.



Conclusion

1. **Big Data & Machine Learning Integration:** Leveraged Apache Spark and Hadoop to process over 44,000 images and metadata, enabling scalable data storage and computation.
2. **Actionable Insights:** Identified patterns in pricing, brand popularity, and demographic trends, aiding retailers in strategic decision-making.
3. **Machine Learning Success:** Applied Linear Regression with high predictive accuracy for inventory optimization and marketing enhancement.
4. **Industry Transformation:** Demonstrated the potential of big data and ML to enhance decision-making, customer experiences, and competitiveness in fashion retail.
5. **Future Scope:** Suggests incorporating real-time data streams and advanced predictive models to refine trend forecasting and customer segmentation.

Contribution

Avirit Singh- Insights, Report , ML model, Data Pre-Processing

Jay Joshi- Insights, Demo, Data Handling

Tanu Datt- Insights, Report , Data Pre-Processing

Vaibhavi Rao- EDA , Insights, PPT

Varun Patil- Insights, Report, PPT, ML Model

Pragya Priyadarshini- Insights , PPT, Data Pre-Processing